CHANDRA KISHORE REDDY GURRAM

Open to Relocation • (623) 286-7971 • gurramchandrakishorereddy@gmail.com • [LinkedIn](#) • [GitHub](#)

Software Engineer with 1 year of experience building scalable systems and intelligent backend services. Skilled in LangChain, OpenAI, and vector DBs (PGVector, PostgreSQL), with hands-on work in document parsing, chunking, and semantic retrieval. Proven track record delivering data-driven features, CI/CD automation, and cloud-native deployments (AWS, Docker, Kubernetes).

EDUCATION

Arizona State University, Tempe, AZ, USA | Masters in Computer Science, (GPA: 4.0/4.0)   Aug 2023 – May 2025
VIT | Integrated Masters in Software Engineering, (GPA: 8.4/10.0)

Aug 2017 – May 2022

TECHNICAL SKILLS

Languages: Python, JavaScript, Java, SQL, C#, C, C++, MATLAB, HTML, CSS, Bash
Frameworks: Django, Flask, Spring Boot, Maven, Hibernate, JUnit, TestNG, React.js, Node.js, Vue.js, Angular.js, Express.js, D3.js, Next.js, Axios, Windows Forms, ADO.NET, ASP.NET, Bootstrap, Robot Framework
Tools And Technologies: AWS SQS, AWS Lambda, AWS ECR, AWS S3, MySQL, PostgreSQL, MongoDB, Hadoop, Spark, Neo4j, Kafka, Kibana, Docker, Git, Kubernetes, Jenkins, JIRA, Artifactory, Selenium
Software Development Practices & Protocols: CI/CD, SDLC, OOPS, REST, RPC, GraphQL, NETCONF, Agile, Distributed Systems, Software Design Patterns, Microservices Architecture, Code Reviews
Certification: [AWS Certified Solutions Architect Associate](#)

PROFESSIONAL EXPERIENCE

Software Engineer Intern   June 2025 – August 2025
Karibu.ai

- **Developed the AI agent** behind Karibu, a context-aware onboarding assistant that guides healthcare professionals in onboarding.
- Designed and implemented a **RAG system** using LangChain to enable semantic search over customer knowledge bases.

Software Development Engineer   May 2022 – July 2023
Nokia Networks and Solutions

- Architected **distributed RESTful APIs** with Flask and MongoDB, resolving NETCONF caching latency for real-time config access.
- Developed CI/CD-integrated framework for config file **backward compatibility, automating validation** to enable higher-quality releases.
- **Implemented lazy loading** for topology data in React, restructuring payload flow to eliminate UI timeouts and deliver instant graph views.
- **Improved CI/CD workflows** by fixing test failures, identifying root causes, and eliminating recurring issues.
- Designed and implemented from scratch a **configuration versioning and monitoring** system enabling full change history, version comparison, rollback, and real-time device metric correlation.
- Developed complete backend architecture including database schema, RESTful APIs, validation pipeline, and live dashboard integration for end-to-end **configuration lifecycle managemen**t.

Student Intern   Sep 2021 – May 2022
Nokia Networks and Solutions

- Automated Controller setup by **scripting Kubernetes pod provisioning** in Python, resolving inconsistency and saving manual hours.
- Implemented a test **execution framework**, **enabling focused local automation runs**, addressing feedback loops and improved team efficiency.
- Built a **Python artifact validation workflow** to parse release notes and flag broken links, addressing error-prone manual checks in releases.

ACADEMIC PROJECTS

Data Engineering and Analytics (Apache Spark, SparkSQL, Scala, GeoSpark, Neo4j)
- Orchestrated a streaming data pipeline with Kubernetes, Kafka, and Neo4j to power real-time analytics for high-volume Reddit datasets.
- Accelerated analytics by leveraging pg bulkload, SparkSQL spatial queries, and PostgreSQL partitioning, ensuring rapid data accessibility.

Exploratory Data Analysis (React, D3.js, React.js, JavaScript, Node.js)
- Engineered a city analytics dashboard via React and D3.js, delivering real-time, SQL-driven visualizations through robust REST APIs.
- Enhanced user insights with interactive UI components and Node.js back end, supporting live data exploration and dynamic analytics.

Generative AI: ASU's Text-To-Video Model
- Curated large-scale video datasets (Panda70M, OpenVid, VidGen) to fuel ASU's text-to-video generation pipeline for academic research.
- Benchmarked LLaVA, QWEN, ORYX using EMScore, SPICE, BLEU, iterating prompt designs to elevate caption accuracy and output.

Multi-Tier Elastic Web Application (AWS, EC2, S3, SQS, RDS, ELB, CloudWatch, AWS Lambda)
- Delivered a scalable face recognition application on AWS with Spring Boot, integrating EC2, S3, RDS, SQS, and elastic load balancing.
- Automated auto-scaling and zero-downtime video processing through SQS, CloudWatch, AWS Lambda, and Spring Cloud orchestration.