

REDDY GURRAM

Tempe, AZ (Open to Relocation) • 480-847-0240 • gurramchandrakishoreddy@gmail.com • [LinkedIn](#) • [GitHub](#)

- Software Engineer with around **2+ years** of experience building fullstack projects and AI-powered software tools.
- Skilled in Core Java, SpringBoot, Django, AI Agent development and RAG pipelines.
- Proven track record delivering data-driven features, CI/CD automation, and cloud-native deployments (AWS, Vercel, Kubernetes).

EDUCATION

Arizona State University, Tempe, AZ, USA | Masters in Computer Science, (GPA: 4.0/4.0)
VIT | Integrated Masters in Software Engineering, (GPA: 8.4/10.0)

Aug 2023 – May 2025
Aug 2017 – May 2022

TECHNICAL SKILLS

Languages: Java, Python, JavaScript, SQL

Frameworks: Spring Boot, Hibernate, JUnit, Django, Flask, FastAPI, Maven, React.js, Node.js, Express.js, D3.js, Next.js.

Tools And Technologies: AWS, PostgreSQL, MySQL, MongoDB, Hadoop, Spark, Neo4j, Kafka, Docker, Git, Kubernetes, Jenkins, Jira.

Software Development Practices & Protocols: CI/CD, SDLC, OOPS, REST, RPC, GraphQL, NETCONF, Agile, Distributed Systems, Software Design Patterns, Microservices Architecture, Code Reviews

Certification: [AWS Certified Solutions Architect Associate](#)

PROFESSIONAL EXPERIENCE

Software Engineer

ASU Laboratory for Energy And Power Solutions (LEAPS)

November 2025 – Present
Mesa, AZ, USA

- Developing a prototype using **Django MVT** to visualize microgrid layouts, single-line diagrams, agent outputs, and support funding demos.
- Building **multi-agent** systems using MCP servers and the OpenAI Agents SDK to automate microgrid design and optimization from satellite imagery.
- Implementing an **orchestrator** that synchronizes independently running agents into a unified microgrid design and optimization workflow.

Software Engineer

[Karibu.ai](#)

June 2025 – November 2025
San Francisco, CA, USA

- Built a **Django REST Framework backend** with **40+ REST API** endpoints using ViewSets, serializers, JWT authentication, and role-based permissions, integrating PostgreSQL, Redis, OpenAI, and HeyGen services.
- **Designed and owned the end-to-end development** of the DNA Management module in Django, extracting key organizational information and directly contributing to onboarding **12 pilot customers**.
- Built robust **Django ORM models** and backend logic supporting multi-stage workflows, adding **Redis-based locking** to ensure consistent progress tracking and concurrency safety.
- Developed a **RAG retrieval service in Python** using LangChain, OpenAI embeddings, and a pgvector-backed **PostgreSQL** store, exposing it cleanly through **Django/DRF APIs** for contextual data retrieval.
- Built and integrated a **Next.js + React frontend** that consumes **DRF APIs**, delivering conversational interfaces, admin dashboards, and real-time avatar-driven user experiences.

Software Development Engineer

Nokia Networks and Solutions

August 2021 – August 2023
Chennai, India

- Built **Spring Boot (Core Java) APIs** for config backward compatibility using **Streams** and custom validators, integrating with Jenkins.
- Implemented lazy loading and optimized **DTO serialization** with **Java Collections** and caching, eliminating topology UI timeouts.
- Improved **CI/CD stability** by diagnosing **Java microservice** failures via **JVM-level debugging** and resolving recurring pipeline issues.
- Developed config versioning and monitoring features using **Spring Boot and Node.js with PostgreSQL** for consistent controller state tracking.
- Automated Controller setup using Python orchestration and Java utilities, streamlining **Kubernetes pod provisioning** and reducing setup time by 70%.
- Built a **lightweight Core Java + JUnit automation framework** with Executors/Futures to speed up local test runs and shorten feedback loops.
- Created an **artifact validation** workflow by combining **Python parsers with Java tooling** to detect release-note issues and replace manual checks.

ACADEMIC PROJECTS

Data Engineering and Analytics (Apache Spark, SparkSQL, Scala, GeoSpark, Neo4j, Kubernetes, Docker, Kafka, Neo4j)

- Orchestrated a streaming **data pipeline with Kubernetes, Kafka, and Neo4j** to power real-time analytics for high-volume Reddit datasets.
- Accelerated analytics by leveraging pg bulkload, SparkSQL spatial queries, and PostgreSQL partitioning, ensuring rapid data accessibility.

Exploratory Data Analysis (React, D3.js, React.js, JavaScript, Node.js)

- Engineered a city analytics dashboard via **React and D3.js**, delivering real-time, SQL-driven visualizations through robust **REST APIs**.
- Enhanced user insights with interactive UI components and Node.js back end, supporting live data exploration and dynamic analytics.

Multi-Tier Elastic Web Application (AWS, EC2, S3, SQS, RDS, ELB, CloudWatch, AWS Lambda)

- Delivered a scalable face recognition application on AWS with Spring Boot, integrating EC2, S3, RDS, SQS, and elastic load balancing.
- Automated auto-scaling and zero-downtime video processing through SQS, CloudWatch, AWS Lambda, and Spring Cloud orchestration.

Rich Text Doc (PDF) RAG for Arxiv Papers (Apache Kafka, RocksDB, Faiss, S3, Lambda, OpenAI)

- Constructed an **RAG pipeline for arXiv papers** using S3, Lambda, RocksDB, FAISS, and Kafka to streamline LLM embedding retrieval.
- Enabled rapid and scalable document search by integrating FAISS indexing with Kafka-based pipelines, elevating query performance.